

# Ray Tracing Acceleration

CS 4620 Lecture 22


# Topics

- Transformations in ray tracing
  - Transforming objects
  - Transformation hierarchies
- Ray tracing acceleration structures
  - Bounding volumes
  - Bounding volume hierarchies
  - Uniform spatial subdivision
  - Adaptive spatial subdivision

# Transforming objects

- In modeling, we've seen the usefulness of transformations
  - How to do the same in RT?
- Take spheres as an example: want to support transformed spheres
  - Need a new Surface subclass
- Option 1: transform sphere into world coordinates
  - Write code to intersect arbitrary ellipsoids
- Option 2: transform ray into sphere's coordinates
  - Then just use existing sphere intersection routine

# Intersecting transformed objects



# Implementing RT transforms

- Create wrapper object “TransformedSurface”
  - Subclass of Surface
  - Has a transform  $T$  and a reference to a surface  $S$
  - To intersect:
    - Transform ray to local coords (by inverse of  $T$ )
    - Call `surface.intersect`
    - Transform hit data back to global coords (by  $T$ )
      - Intersection point
      - Surface normal
      - Any other relevant data (maybe none)

```
class TransformedSurface : Surface {  
    Transform xf;  
    Surface s;  
    intersect(Ray r, tMin, tMax) {  
        (p, n, t) = s.intersect(xf-1 r, tMin, tMax);  
        return (xf p, xf-T n, t);  
    }  
}
```

# Groups, transforms, hierarchies

- Often it's useful to transform several objects at once
  - Add “SurfaceGroup” as a subclass of Surface
    - Has a list of surfaces
    - Returns closest intersection
      - Opportunity to make Scene a single Surface to avoid duplication
- With TransformedSurface and SurfaceGroup you can put transforms below transforms
  - Voilà! A transformation hierarchy.

```

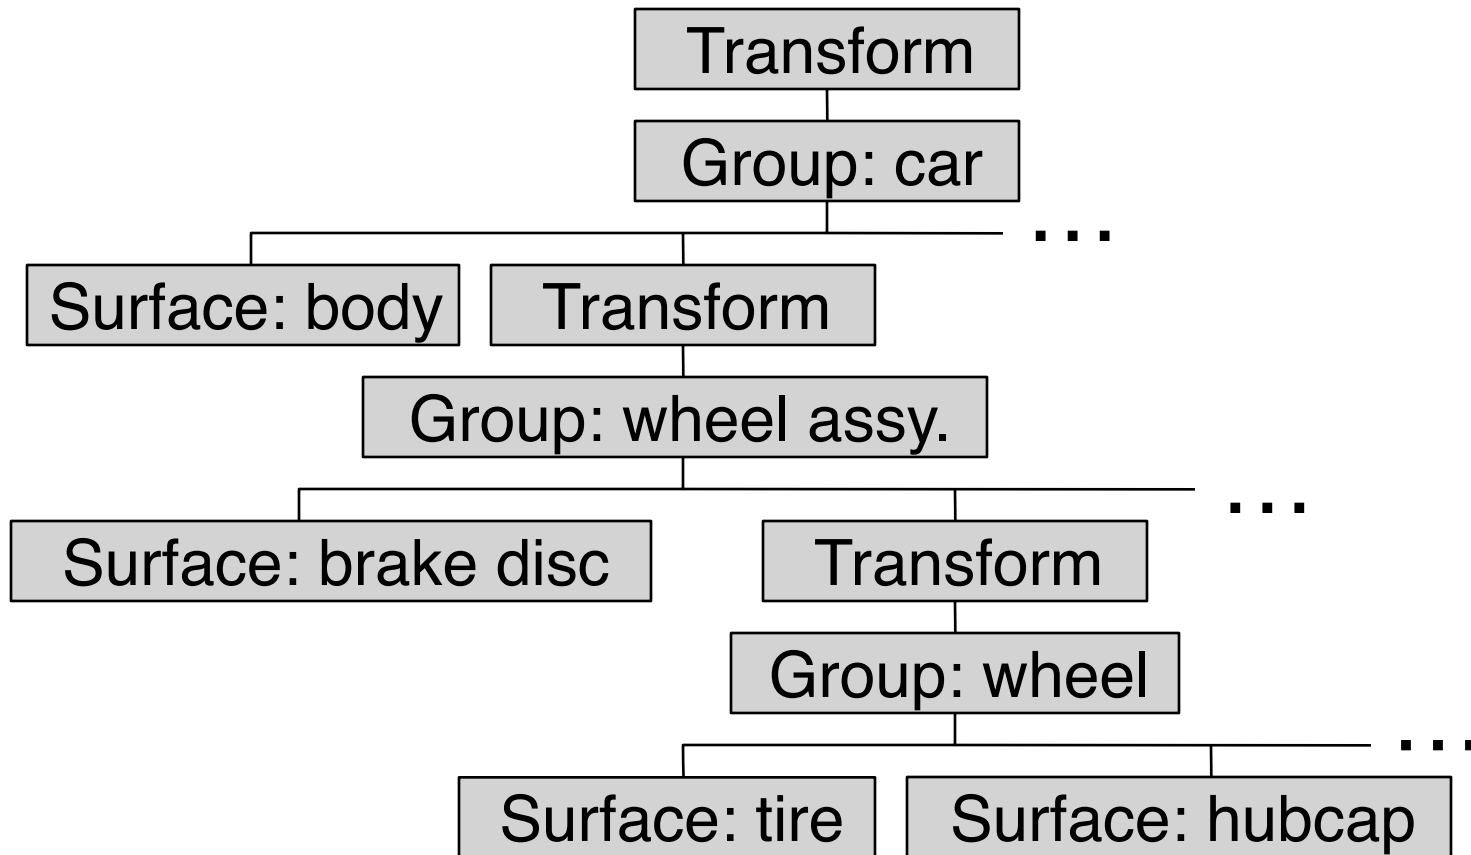
class SurfaceGroup : Surface {
    List<Surface> sList;
    intersect(Ray r, tMin, tMax) {
        hit = false;
        for (s in sList) {
            result = s.intersect(r, tMin, tMax);
            if (result) {
                (p,n,t) = result
                tMax = min(t, tMax);
                hit = true;
            }
        }
        if (hit)
            return (p, n, tMax);
        else
            return false;
    }
}

```

Algorithm is quite familiar from Scene.intersect in Ray 1...



# A transformation hierarchy

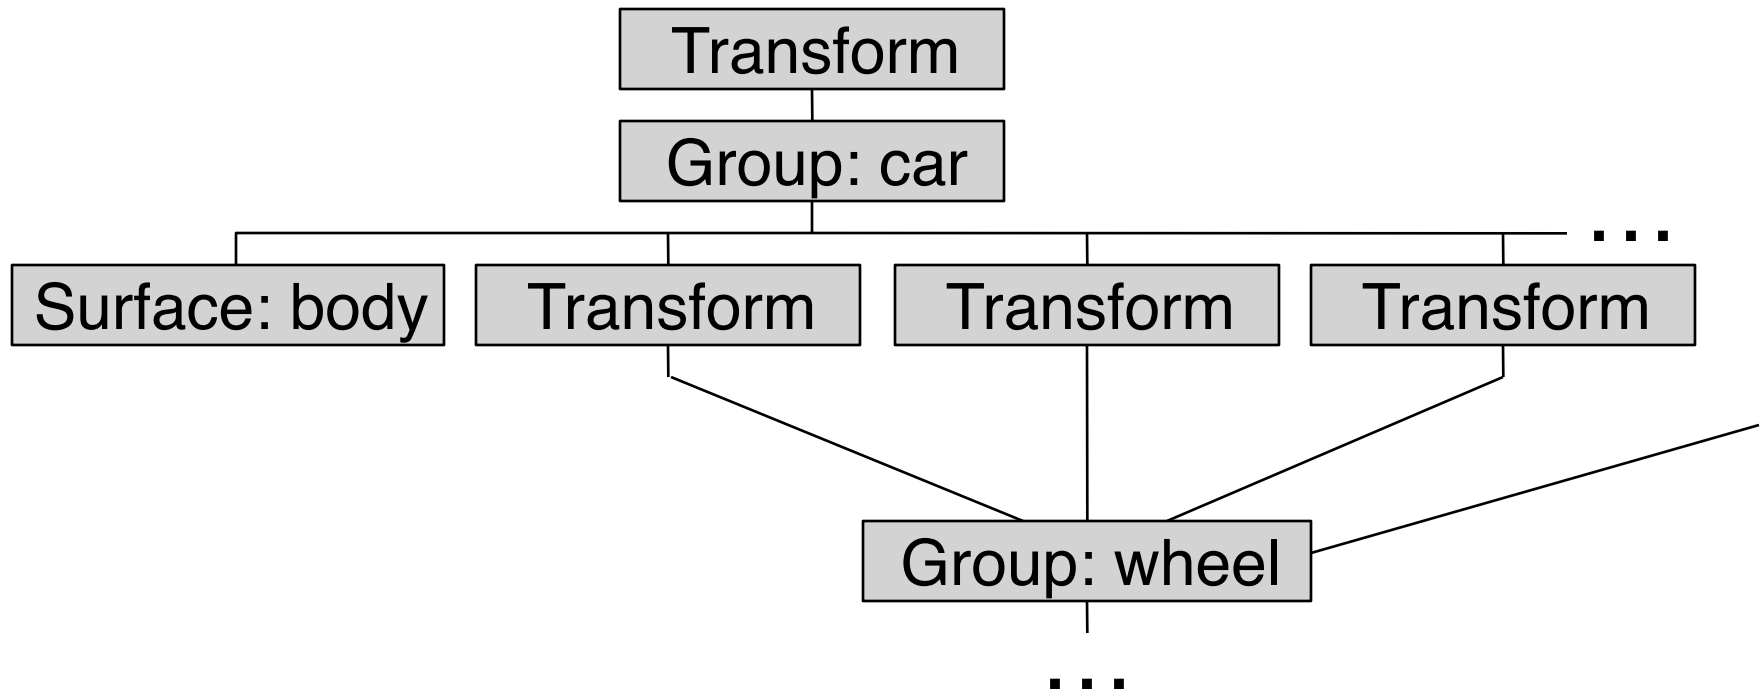


- Common optimization: merge transforms with groups

# Instancing

- Anything worth doing is worth doing  $n$  times
- If we can transform objects, why not transform them several ways?
  - Many models have repeated subassemblies
    - Mechanical parts (wheels of car)
    - Multiple objects (chairs in classroom, ...)
  - Nothing stops you from creating two TransformedSurface objects that reference the same Surface
    - Allowing this makes the transformation tree into a DAG
      - (directed acyclic graph)
    - Mostly this is transparent to the renderer

# Hierarchy with instancing



- Previous code still works just fine!

# Hierarchies and performance

- Transforming rays is expensive
  - minimize tree depth: flatten on input
    - push all transformations toward leaves
    - optional for triangle meshes
      - transform ray once, amortize cost over many intersections
  - internal group nodes still required for instancing
    - can't push two transforms down to same child!

```

TransformGroup {
  xf: A
  Mesh {
    v1, v2, v3, ...
  }
  TransformGroup {
    xf: B
    Sphere {
      radius: r
    }
  }
}

```



```

Mesh {
  xf: A
  v1, v2, v3, ...
}
Sphere {
  xf: AB
  radius: r
}

```



```

Mesh {
  Av1, Av2, Av3, ...
}
Sphere {
  xf: AB
  radius: r
}




```

# Ray tracing acceleration

- Ray tracing is slow. This is bad!
  - Ray tracers spend most of their time in ray-surface intersection methods
- Ways to improve speed
  - Make intersection methods more efficient
    - Yes, good idea. But only gets you so far
  - Call intersection methods fewer times
    - Intersecting every ray with every object is wasteful
    - Basic strategy: efficiently find big chunks of geometry that definitely do not intersect a ray

# Bounding volumes

- Quick way to avoid intersections: bound object with a simple volume
  - Object is fully contained in the volume
  - If it doesn't hit the volume, it doesn't hit the object
  - So test bvol first, then test object if it hits



# Bounding volumes

- Cost: more for hits and near misses, less for far misses
- Worth doing? It depends:
  - Cost of bvol intersection test should be small
    - Therefore use simple shapes (spheres, boxes, ...)
  - Cost of object intersect test should be large
    - Bvols most useful for complex objects
  - Tightness of fit should be good
    - Loose fit leads to extra object intersections
    - Tradeoff between tightness and bvol intersection cost



# Implementing bounding volume

- Just add new Surface subclass, “BoundedSurface”
  - Contains a bounding volume and a reference to a surface
  - Intersection method:
    - Intersect with bvol, return false for miss
    - Return `surface.intersect(ray)`
  - Like transformations, common to merge with group
  - This change is transparent to the renderer (only it might run faster)
- Note that all Surfaces will need to be able to supply bounding volumes for themselves

```
class BoundedSurface : Surface {  
    BVol v;  
    Surface s;  
    intersect(Ray r, tMin, tMax) {  
        if (v.intersect(r, tMin, tMax))  
            return s.intersect(r, tMin, tMax);  
        else  
            return false;  
    }  
}
```


# If it's worth doing, it's worth doing hierarchically!

- Bvols around objects may help
- Bvols around groups of objects will help
- Bvols around parts of complex objects will help
- Leads to the idea of using bounding volumes all the way from the whole scene down to groups of a few objects


# Implementing a bvol hierarchy

- A BoundedSurface can contain a list of Surfaces
- Some of those Surfaces might be more BoundedSurfaces
- Voilà! A bounding volume hierarchy
  - And it's all still transparent to the renderer


# BVH construction example




# BVH construction example



# BVH construction example




# BVH construction example







# BVH ray-tracing example




# BVH ray-tracing example




# BVH ray-tracing example




# BVH ray-tracing example




# BVH ray-tracing example




# BVH ray-tracing example



# BVH ray-tracing example




# BVH ray-tracing example






# BVH ray-tracing example



# BVH ray-tracing example





# Choice of bounding volumes

- Spheres -- easy to intersect, not always so tight
- Axis-aligned bounding boxes (AABBs) -- easy to intersect, often tighter (esp. for axis-aligned models)
- Oriented bounding boxes (OBBs) -- easy to intersect (but cost of transformation), tighter for arbitrary objects
- Computing the bvols
  - For primitives -- generally pretty easy
  - For groups -- not so easy for OBBs (to do well)
  - For transformed surfaces -- not so easy for spheres

# Axis aligned bounding boxes

- Probably easiest to implement
- Computing for (axis-aligned) primitives
  - Cube: duh!
  - Sphere, cylinder, etc.: pretty obvious
  - Groups or meshes: min/max of component parts
- AABBs for transformed surface
  - Easy to do conservatively: bbox of the 8 corners of the bbox of the untransformed surface
- How to intersect them
  - Treat them as an intersection of slabs (see Shirley)

# Intersecting boxes




$$t \in [t_{xmin}, t_{xmax}]$$

$$t \in [t_{ymin}, t_{ymax}]$$

$$t \in [t_{xmin}, t_{xmax}] \cap [t_{ymin}, t_{ymax}]$$


# Ray-box intersection

- Could intersect with 6 faces individually
- Better way: box is the intersection of 3 slabs




# Ray-box intersection

- Could intersect with 6 faces individually
- Better way: box is the intersection of 3 slabs



# Ray-box intersection


- Could intersect with 6 faces individually
- Better way: box is the intersection of 3 slabs





# Ray-box intersection

- Could intersect with 6 faces individually
- Better way: box is the intersection of 3 slabs



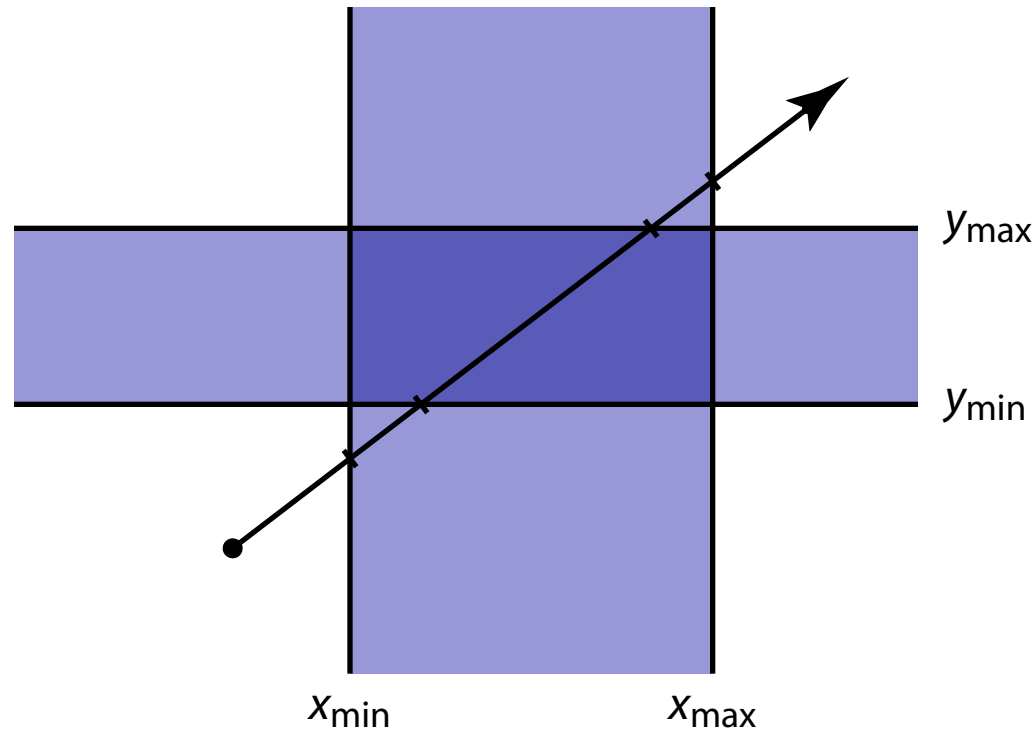
# Ray-slab intersection

- 2D example
- 3D is the same!



# Ray-slab intersection

- 2D example
- 3D is the same!

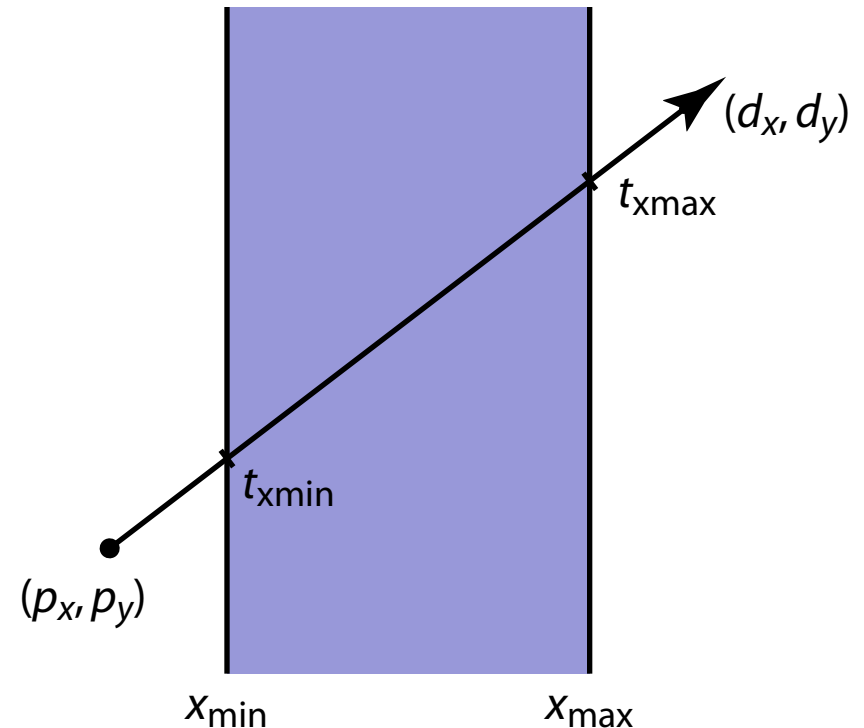


# Ray-slab intersection

- 2D example
- 3D is the same!

$$p_x + t_{x\min} d_x = x_{\min}$$

$$t_{x\min} = (x_{\min} - p_x) / d_x$$



# Ray-slab intersection


- 2D example
- 3D is the same!

$$p_x + t_{x\min} d_x = x_{\min}$$

$$t_{x\min} = (x_{\min} - p_x) / d_x$$

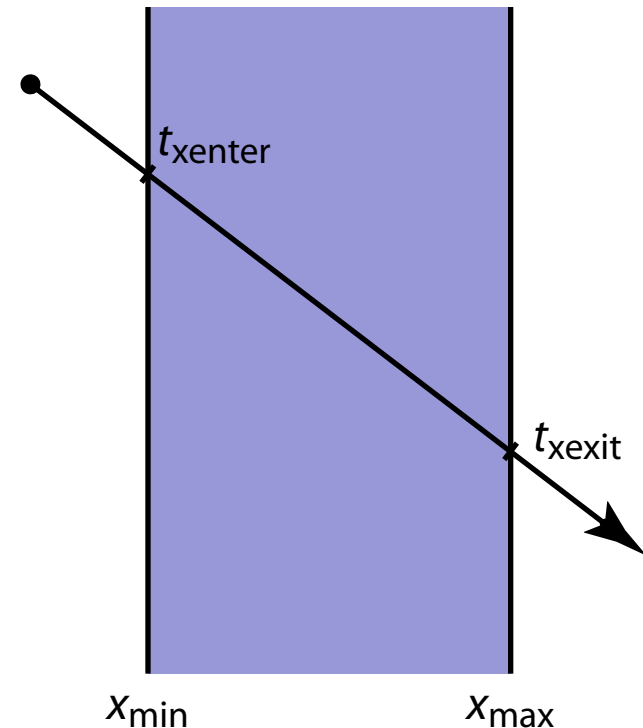
$$p_y + t_{y\min} d_y = y_{\min}$$

$$t_{y\min} = (y_{\min} - p_y) / d_y$$



# Intersecting intersections

- Each intersection is an interval
- Want last entry point and first exit point

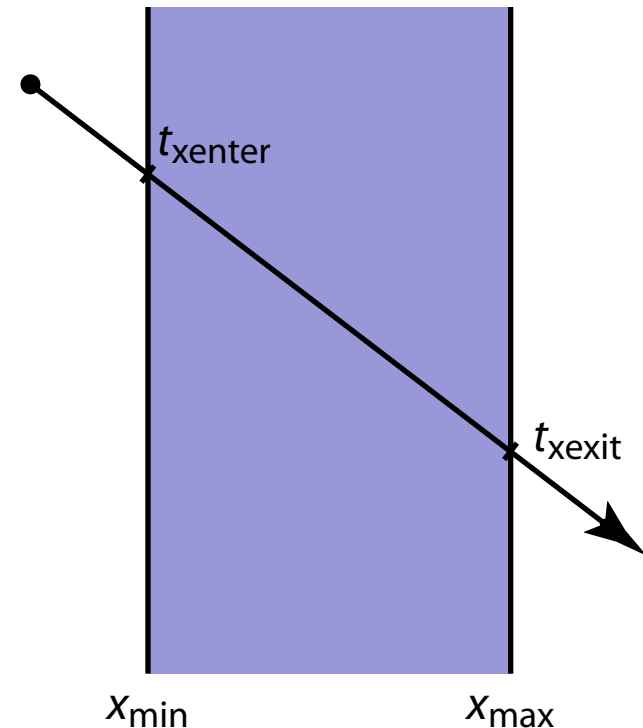


# Intersecting intersections

- Each intersection is an interval
- Want last entry point and first exit point

$$t_{xenter} = \min(t_{xmin}, t_{xmax})$$

$$t_{xexit} = \max(t_{xmin}, t_{xmax})$$




# Intersecting intersections

- Each intersection is an interval
- Want last entry point and first exit point

$$t_{xenter} = \min(t_{xmin}, t_{xmax})$$

$$t_{xexit} = \max(t_{xmin}, t_{xmax})$$





# Intersecting intersections

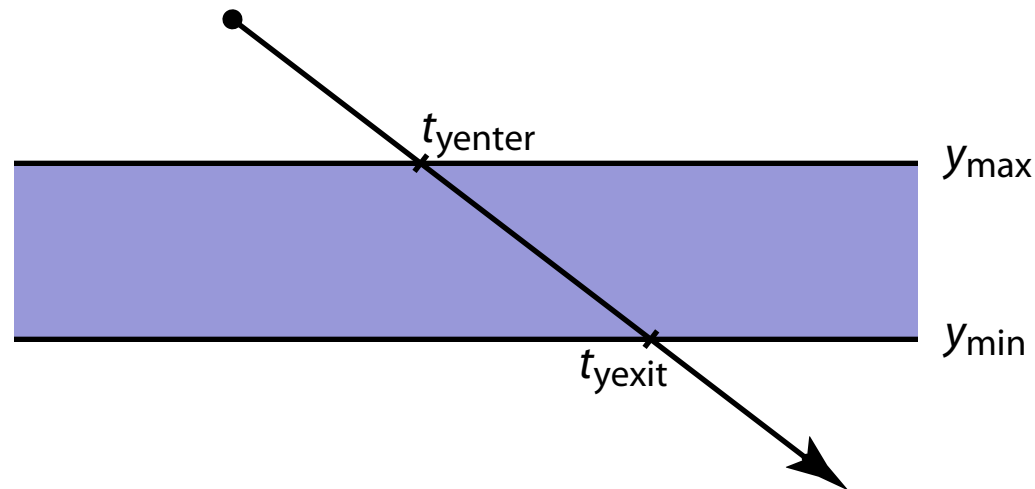
- Each intersection is an interval
- Want last entry point and first exit point

$$t_{xenter} = \min(t_{xmin}, t_{xmax})$$

$$t_{xexit} = \max(t_{xmin}, t_{xmax})$$

$$t_{yenter} = \min(t_{ymin}, t_{ymax})$$

$$t_{yexit} = \max(t_{ymin}, t_{ymax})$$



# Intersecting intersections

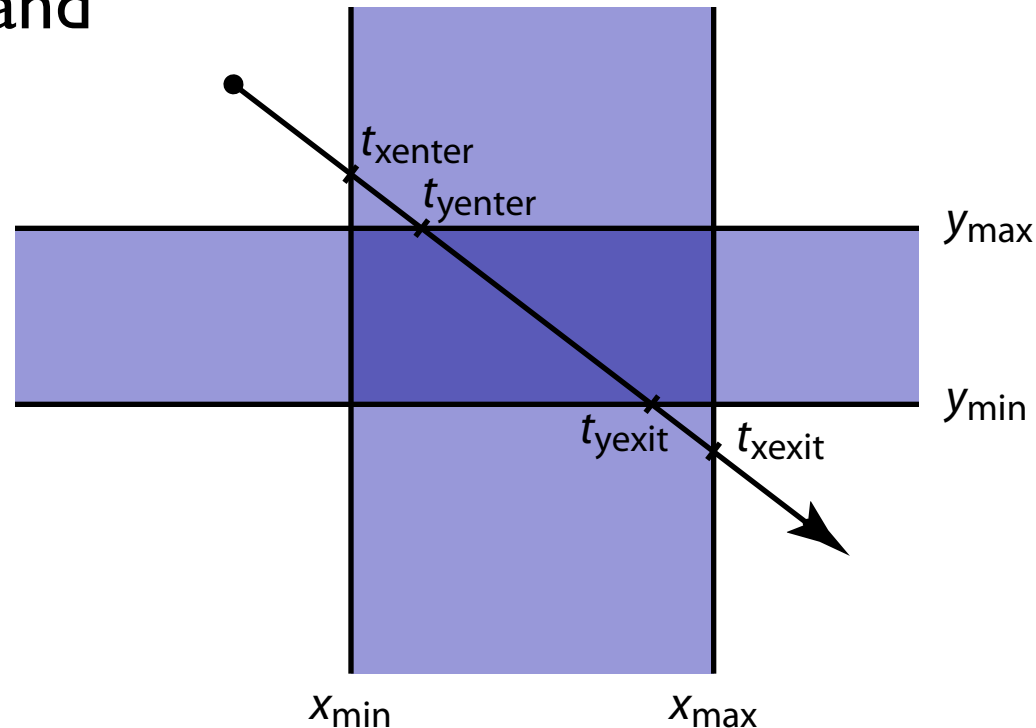
- Each intersection is an interval
- Want last entry point and first exit point

$$t_{xenter} = \min(t_{xmin}, t_{xmax})$$

$$t_{xexit} = \max(t_{xmin}, t_{xmax})$$

$$t_{yenter} = \min(t_{ymin}, t_{ymax})$$

$$t_{yexit} = \max(t_{ymin}, t_{ymax})$$



# Intersecting intersections

- Each intersection is an interval
- Want last entry point and first exit point

$$t_{xenter} = \min(t_{xmin}, t_{xmax})$$

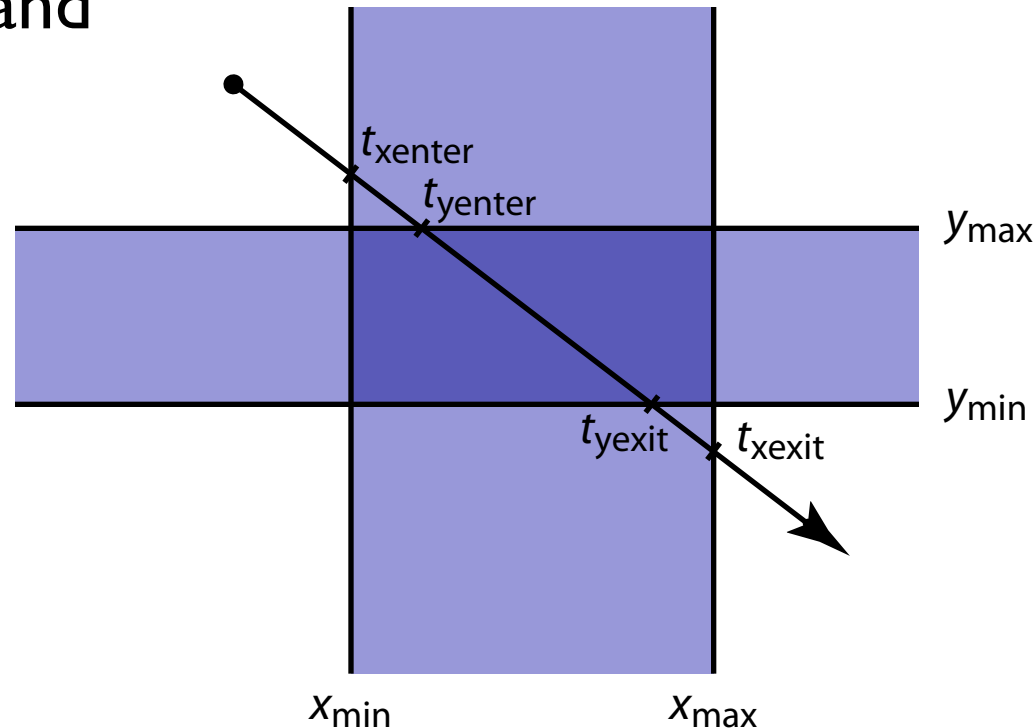
$$t_{xexit} = \max(t_{xmin}, t_{xmax})$$

$$t_{yenter} = \min(t_{ymin}, t_{ymax})$$

$$t_{yexit} = \max(t_{ymin}, t_{ymax})$$

$$t_{enter} = \max(t_{xenter}, t_{yenter})$$

$$t_{exit} = \min(t_{xexit}, t_{yexit})$$



# Building a hierarchy


- Usually do it top-down
- Make bbox for whole scene, then split into (maybe 2) parts
  - Recurse on parts
  - Stop when there are just a few objects in your box

# Building a hierarchy

- How to partition?
  - Ideal: clusters
  - Practical: partition along axis
    - Center partition
      - Less expensive, simpler
      - Unbalanced tree (but may sometimes be better)
    - Median partition
      - More expensive
      - More balanced tree
    - Surface area heuristic
      - Model expected cost of ray intersection
      - Generally produces best-performing trees

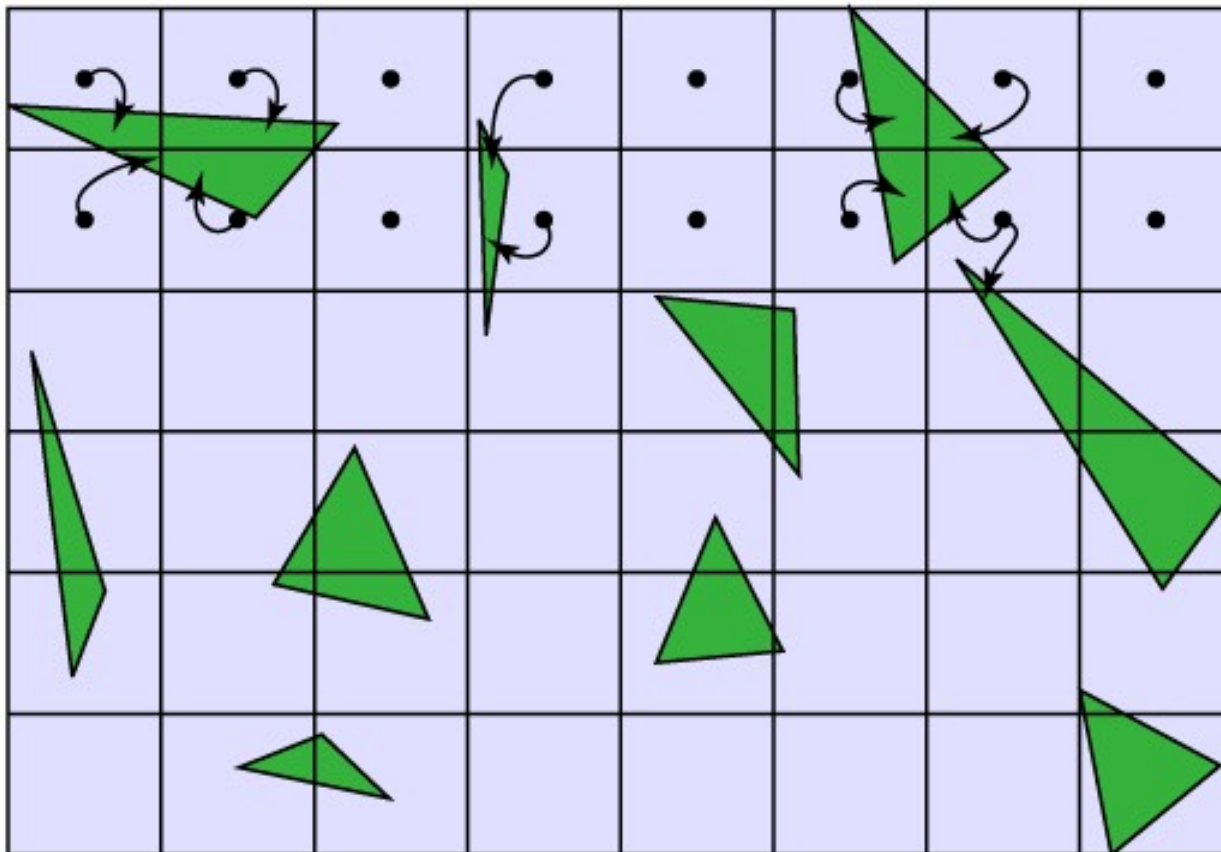
# Regular space subdivision

- An entirely different approach: uniform grid of cells

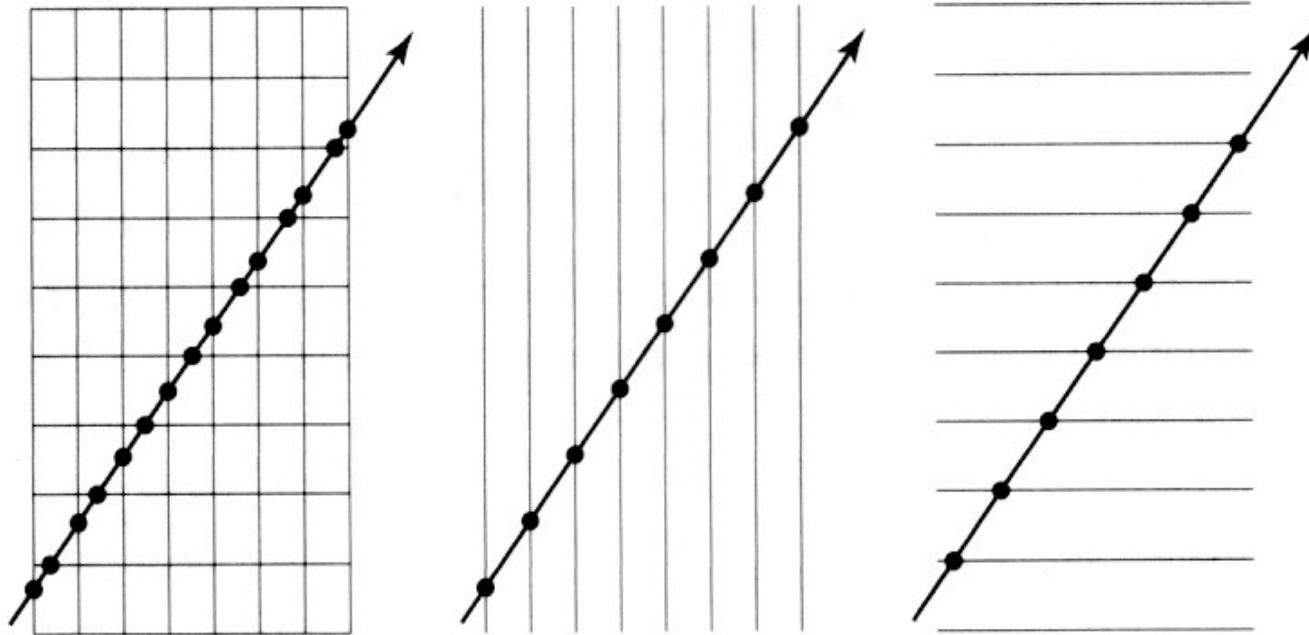


# Regular grid example

- Grid divides space, not objects




# Traversing a regular grid






# Non-regular space subdivision

- *k*-d Tree
  - subdivides space, like grid
  - adaptive, like BVH




# Non-regular space subdivision

- *k*-d Tree
  - subdivides space, like grid
  - adaptive, like BVH



# Non-regular space subdivision

- *k*-d Tree
  - subdivides space, like grid
  - adaptive, like BVH



# Implementing acceleration structures

- Conceptually simple to build acceleration structure into scene structure
  - what we sketched out earlier
- Better engineering decision to separate them
  - plug and play different acceleration structures
  - keep representation of scene itself simple